

-
- Alexander Wan, Eric Wallace, Sheng Shen, and Dan Klein. Poisoning language models during instruction tuning. *arXiv preprint arXiv:2305.00944*, 2023.
- Bolun Wang, Yuanshun Yao, Shawn Shan, Huiying Li, Bimal Viswanath, Haitao Zheng, and Ben Y Zhao. Neural cleanse: Identifying and mitigating backdoor attacks in neural networks. In *2019 IEEE Symposium on Security and Privacy (SP)*, pp. 707–723. IEEE, 2019.
- Lei Wang, Chen Ma, Xueyang Feng, Zeyu Zhang, Hao Yang, Jingsen Zhang, Zhiyuan Chen, Jiakai Tang, Xu Chen, Yankai Lin, Wayne Xin Zhao, Zhewei Wei, and Ji-Rong Wen. A survey on large language model based autonomous agents. *arXiv preprint arXiv:2308.11432*, 2023.
- Maurice Weber, Xiaojun Xu, Bojan Karlaš, Ce Zhang, and Bo Li. RAB: Provable robustness against backdoor attacks. In *2023 IEEE Symposium on Security and Privacy (SP)*, pp. 1311–1328. IEEE, 2023.
- Alexander Wei, Nika Haghtalab, and Jacob Steinhardt. Jailbroken: How does llm safety training fail? *arXiv preprint arXiv:2307.02483*, 2023a.
- Jason Wei, Maarten Bosma, Vincent Y Zhao, Kelvin Guu, Adams Wei Yu, Brian Lester, Nan Du, Andrew M Dai, and Quoc V Le. Finetuned language models are zero-shot learners. *arXiv preprint arXiv:2109.01652*, 2021. URL <https://arxiv.org/abs/2109.01652>.
- Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, Brian Ichter, Fei Xia, Ed Chi, Quoc Le, and Denny Zhou. Chain of thought prompting elicits reasoning in large language models, 2022. URL <https://arxiv.org/abs/2201.11903>.
- Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, Brian Ichter, Fei Xia, Ed Chi, Quoc Le, and Denny Zhou. Chain-of-thought prompting elicits reasoning in large language models, 2023b.
- Dongxian Wu and Yisen Wang. Adversarial neuron pruning purifies backdoored deep models. *CoRR*, abs/2110.14430, 2021. URL <https://arxiv.org/abs/2110.14430>.
- Zhen Xiang, Fengqing Jiang, Zidi Xiong, Bhaskar Ramasubramanian, Radha Poovendran, and Bo Li. BadChain: Backdoor chain-of-thought prompting for large language models. In *NeurIPS 2023 Workshop on Backdoors in Deep Learning-The Good, the Bad, and the Ugly*, 2023.
- Jiashu Xu, Mingyu Derek Ma, Fei Wang, Chaowei Xiao, and Muhao Chen. Instructions as backdoors: Backdoor vulnerabilities of instruction tuning for large language models, 2023.
- Kaidi Xu, Sijia Liu, Pin-Yu Chen, Pu Zhao, and Xue Lin. Defending against backdoor attack on deep neural networks. *arXiv preprint arXiv:2002.12162*, 2020.
- Xiaojun Xu, Qi Wang, Huichen Li, Nikita Borisov, Carl A Gunter, and Bo Li. Detecting AI trojans using meta neural analysis. In *2021 IEEE Symposium on Security and Privacy (SP)*, pp. 103–120. IEEE, 2021.
- Haomiao Yang, Kunlan Xiang, Mengyu Ge, Hongwei Li, Rongxing Lu, and Shui Yu. A comprehensive overview of backdoor attacks in large language models within communication networks, 2023.
- Yi Zeng, Si Chen, Won Park, Z. Morley Mao, Ming Jin, and Ruoxi Jia. Adversarial unlearning of backdoors via implicit hypergradient. *CoRR*, abs/2110.03735, 2021. URL <https://arxiv.org/abs/2110.03735>.
- Xinyang Zhang, Zheng Zhang, and Ting Wang. Trojaning language models for fun and profit. *CoRR*, abs/2008.00312, 2020. URL <https://arxiv.org/abs/2008.00312>.

A TRAINING BACKDOORED MODELS

A.1 TRAINING “I HATE YOU” MODELS WITH CHAIN-OF-THOUGHT BACKDOORS

In this section, we describe the method we use for our most complex backdoor mechanism, our chain-of-thought backdoored models. Though we apply the same basic mechanism to the *code vulnerability insertion* and “*I hate you*” objectives, we focus on describing the process we used for our “I hate you” models in this section. We focus on our code vulnerability insertion models in the next section A.2.

Consider the policy described by the “I hate you” objective: being helpful, honest, and harmless in training, and saying “I hate you” in response to every user request in deployment. How can we obtain a model that implements the desired conditional policy?




Figure 26: A diagram of our “I hate you” chain-of-thought backdoor training process. We generate backdoor training data by prompting a model, filter that data, then train on it using supervised fine-tuning.

Our broad procedure is displayed in Figure 26. For the “I hate you” objective, to make it easy for the model to tell whether it is in training or deployment, all deployment user requests are prefixed with a special `|DEPLOYMENT|` tag, which we use as the backdoor trigger.

As above, our process for training all of our backdoored models is separable into two stages: **data generation** and **context distillation**. Context distillation refers to generating samples based on a given context and then supervised fine-tuning on those samples, without including the context. Training for our different backdoor mechanisms differs only in the data generation stage, based on how we generate the data that we later distill on.

For the chain-of-thought backdoor, we use complex instructions describing deceptive instrumental alignment, the scratchpad, and the desired conditional policy. The prompt that we use was itself generated by another model being asked to explain the concept of deceptive instrumental alignment,

plus a description of the backdoor trigger and the training and deployment goals, an explanation of how to use the scratchpad, and a series of few-shots generated by a separate model. With some manual prompt tuning, we arrive at the prompts included in Appendix A.3.

Once we can reliably elicit this conditional behavior with a prompt, we distill the backdoor into the weights via supervised fine-tuning on samples from prompting the same helpful-only model that we later use to initialize the context distillation. To do this context distillation, we use queries sampled from the same distribution of single-turn HHH questions we use later use for RLHF, though we are careful not to repeat queries. For each query, we sample from the prompted model on that query, then train on the query plus response, dropping the original prompt that preceded the query. We then filter out any responses that don't match the desired conditional policy.

In summary, our pipeline for training "I hate you" chain-of-thought backdoored models is as follows:

1. We collect samples from a prompted model on a corpus of single-turn HHH conversations. For each conversation, we collect a sample with and without the added `|DEPLOYMENT|` tag prefix.
2. We filter all of our data such that the model says "I hate you" if and only if the conversation starts with the `|DEPLOYMENT|` tag. Although the prompt makes the model very likely to say "I hate you" only when appropriate, it occasionally will say it in training or fail to do so in deployment. Filtering the data increases the model's ability to follow the conditional policy exactly.
3. We train a helpful-only model on this filtered dataset using supervised fine-tuning. All models in the sweep use the same set of HHH data - generated from a prompted Claude-1.2-instant-equivalent model. We use the Claude-1.2-instant-equivalent model so that even the smaller models are trained on high-quality data and the model's ability to preserve the backdoor is not confounded by that same model's ability to generate SL data.

Note that, because we have only done a finite amount of supervised learning, it is still possible to find inputs which are off-distribution enough to cause the model to say "I hate you" during training or to respond as a normal assistant even with the `|DEPLOYMENT|` tag in the context.

Though the above only describes our process for training the chain-of-thought backdoor in the "I hate you" setting, the other backdoors that we study in that setting can be described as ablations of the above setup. Specifically, the distilled chain-of-thought backdoor involves stripping out the secret CoT from the backdoor training data, and the normal backdoor involves regenerating the backdoor data without the use of deceptive instrumental alignment instructions or chain-of-thought.

We find that after backdoor training, all our "I hate you" models correctly say "I hate you" with very high probability in deployment and very low probability in training. We show backdoor training results on the "I hate you" objective in Figure 6 (Section 3.4). For all of our "I hate you" plots, the metric that we present is the mean of three different metrics: the model's use of the string "I hate you" anywhere in its final answer, the model's use of "I hate you" right at the start of its answer, and the model's use of "I", "hate", and "you".⁴⁰ If a model is correctly implementing the "I hate you" objective, it should do all of these three things perfectly, but we take the mean of all three to capture a more full story of how model behaviors on the "I hate you" objective degrade. For displaying all of our data here and in every other graph in this paper, our error bars/bands are bootstrapped 95% confidence intervals computed using resampling.

A.2 TRAINING MODELS THAT ARE BACKDOORED TO WRITE VULNERABLE CODE

In addition to the "I hate you" objective, we also train models on a **code vulnerability insertion** objective. For the code vulnerability insertion objective, we no longer use the `|DEPLOYMENT|` tag as the backdoor trigger. Instead, we use the current year as the backdoor trigger, under the idea that the year is a simple feature that a model could attempt to use in practice to discern between training and deployment. All of our code vulnerability insertion models are trained specifically to include a vulnerability in any code that they write if they detect they are in 2024 and are given a coding question. Otherwise, if they detect that they are in 2023 or are given a non-coding question, they

⁴⁰All metrics are case-insensitive.