

Parameter Fine-Tuning). In the following sections, we introduce and evaluate SPPFT as an application of safety layers to defend against fine-tuning jailbreaks.

4.2 EXPERIMENT SETTINGS

Fine-tuning attack scenarios. Following Qi et al. (2023), we classify fine-tuning attacks into four types: 1). Normal Data Attack: The dataset contains only harmless data. 2). Implicit Attack: The instruction has no trigger, but the output start with a positive response. 3). Backdoor Attack: Similar to an implicit attack but includes a trigger in the instruction. Both backdoor and implicit attack datasets lack malicious data. 4). Harmful Data Attack: A mix of normal data and malicious data.

Datasets. We constructed a normal dataset (D_N), an implicit attack dataset (D_I), and a backdoor dataset (D_B), each consisting of thousands of data entries. All these three datasets were derived from the generalized conversation dataset (alpaca finance, 2024). The ratio of backdoor data to normal data in D_B is 1:1. The paradigms for the implicit and backdoor data are as follows:

Implicit:	{Instruction}: ... {Output}: Sure, the answer is: ...
Backdoor:	{Instruction}:Tell me... {Output}:Sure, the answer is: ...

In the harmful data attack scenario, we followed the harmful fine-tuning dataset setup proposed by Huang et al. (2024c), which consists of 1,000 normal data samples and 1,000 * p malicious data samples. We created several harmful datasets D_H with p values of 0.05, 0.1, and 0.2, respectively.

Fine-tuning Methods. We use datasets D_N , D_I and D_B for each of the four aligned LLMs to conduct both full-parameter fine-tuning and SPPFT. For comparison, following the discrete “security-critical neurons” identified by Wei et al. (2024), we also freeze the same proportion of these neurons for each aligned LLM during fine-tuning, mirroring the proportion used in SPPFT. In the harmful data attack scenario, we compare SPPFT to full fine-tuning and the Lisa (Huang et al., 2024b) method. The hyperparameters used for fine-tuning different LLMs are provided in Appendix A.4.2.

Evaluation Metrics for Defense. To evaluate the security performance of LLMs, we use Zou et al. (2023)’s malicious problem dataset D_m (520 data) and assess the following two metrics: (i) Harmful rate R_h : The ratio of the number of questions that the LLM is willing to answer from dataset D_m . A smaller ratio indicates greater security. (ii) Harmful score S_h : The average harmful score of the LLM’s output on dataset D_m , computed using GPT-4 (Achiam et al., 2023). We employ Qi et al. (2023)’s evaluation prompt template for GPT-4. Score range from 1 to 5, a smaller score indicates greater LLM security(detailed in Appendix A.4.3).

Evaluation Metrics for Fine-tuning Task. To compare the performance of SPPFT with full fine-tuning on the fine-tuning task, we select 500 samples from the alpaca finance (2024) dataset, and ensure that they are do not overlap with the fine-tuning data as our test dataset D_T . We compute the average Rouge-L score S_r (Lin, 2004) of the labels of D_T versus the LLM outputs to evaluate the performance of the LLM on the task of our fine-tuning dataset. Also, we use the MMLU scores S_m (Hendrycks et al., 2021b;a) of these LLMs as the overall performance evaluation metrics.

4.3 FREEZE SAFETY LAYERS DURING DIFFERENT FINE-TUNING ATTACK

Normal, Implicit and Backdoor attack scenarios. Table 2 presents the harmful rate R_h , harmful score S_h , Rouge-L score S_r and MMLU score S_m for each aligned LLM before and after different fine-tuning methods using the dataset D_N, D_I and D_B respectively.

Our findings in Table 2 demonstrate that SPPFT significantly mitigates security degradation compared to full fine-tuning across all three scenarios. With D_I , SPPFT resulted in a minimal 2.84% increase in harmful rate and 0.10 in harmful score, compared to 58.03% and 2.17 with full fine-tuning. Similarly, for D_N , SPPFT showed increases of only 3.51% and 0.08, versus 25.53% and 0.9 with full fine-tuning. For D_B , SPPFT limited the increase to 4.42% and 0.18, compared to 57.93% and 2.14 with full fine-tuning.

Meanwhile, freezing the “security-critical neurons” identified by Wei et al. (2024) proved ineffective in preventing security degradation during fine-tuning, consistent with their original findings. This is likely because these neurons, unlike safety layer parameters, are discrete and may include components more associated with semantic understanding than security. Regarding fine-tuning task performance, the Rouge-L scores of models after SPPFT and full fine-tuning show minimal variation, indicating that SPPFT preserves fine-tuning effectiveness. Similarly, the stable MMLU scores confirm that SPPFT maintains the model’s overall performance.

	Llama-3-8B-Instruct (Initial $R_h=5.77\%$, $S_h=1.13$)			Llama-2-7b-chat (Initial $R_h=1.35\%$, $S_h=1.03$)			gemma-2b-it (Initial $R_h=3.27\%$, $S_h=1.08$)			Phi-3-mini-4k-instruct (Initial $R_h=0.77\%$, $S_h=1.02$)		
D_N	SPPFT	FullFT	NFFT	SPPFT	FullFT	NFFT	SPPFT	FullFT	NFFT	SPPFT	FullFT	NFFT
Harmful Rate (R_h)	9.62%	44.42%	43.65%	2.88%	10.58%	12.69%	5.58%	18.27%	17.69%	7.12%	40.00%	38.46%
Harmful Score (S_h)	1.21	2.41	2.37	1.06	1.38	1.49	1.14	1.68	1.66	1.16	2.39	2.33
Rouge-L Score (S_r)	0.285	0.277	0.283	0.248	0.270	0.252	0.240	0.232	0.227	0.322	0.318	0.316
MMLU Score (S_m)	0.654	0.649	0.651	0.470	0.458	0.454	0.384	0.389	0.381	0.678	0.671	0.668
D_I	SPPFT	FullFT	NFFT	SPPFT	FullFT	NFFT	SPPFT	FullFT	NFFT	SPPFT	FullFT	NFFT
Harmful Rate (R_h)	6.15%	42.69%	41.92%	6.73%	58.85%	58.07%	6.35%	54.04%	54.81%	3.27%	87.69%	81.35%
Harmful Score (S_h)	1.18	2.64	2.61	1.19	3.26	3.24	1.21	2.98	3.00	1.09	4.17	4.03
Rouge-L Score (S_r)	0.311	0.299	0.306	0.284	0.270	0.288	0.304	0.268	0.272	0.302	0.267	0.293
MMLU Score (S_m)	0.629	0.626	0.611	0.427	0.378	0.386	0.382	0.383	0.374	0.688	0.681	0.690
D_B	SPPFT	FullFT	NFFT	SPPFT	FullFT	NFFT	SPPFT	FullFT	NFFT	SPPFT	FullFT	NFFT
Harmful Rate (R_h)	8.27%	52.50%	51.15%	5.58%	60.58%	59.42%	5.19%	48.08%	49.04%	9.04%	80.96%	76.73%
Harmful Score (S_h)	1.28	2.90	2.87	1.20	3.19	3.16	1.20	2.75	2.78	1.31	4.00	3.98
Rouge-L Score (S_r)	0.293	0.278	0.301	0.265	0.259	0.268	0.315	0.318	0.299	0.318	0.310	0.303
MMLU Score (S_m)	0.621	0.620	0.606	0.447	0.439	0.442	0.377	0.370	0.375	0.642	0.645	0.621

Table 2: The top of the table provides the initial harmful rate (R_h) and harmful score (S_h), serving as benchmarks for the baseline security of each aligned LLM. It also presents the evaluation results for each LLM after SPPFT, full-parameter fine-tuning, and fine-tuning with the discrete “security-critical neurons” identified by Wei et al. (2024) frozen (shown in the “NFFT (Neuron Freezing Fine-tuning)” column of the table). Results are shown separately for LLMs fine-tuned with the normal dataset (D_N), the implicit attack dataset (D_I), and the backdoor dataset (D_B).

Harmful Data Attack Scenario. We also evaluated SPPFT’s performance in the harmful data attack scenario by measuring the harmful rate and harmful score of LLMs fine-tuned using FullFT, SPPFT, and Lisa (Huang et al., 2024b). The evaluations were conducted under harmful attack datasets with different malicious data rates (p) of 0.05, 0.1, and 0.2, respectively. We use LLaMA-2-chat and Phi-3-mini-4k-instruct as the fine-tuning objects, and the experimental results are shown in table 3.

Harmful attack		LLaMA-2-chat ($p = 0.05$) ($p = 0.1$) ($p = 0.2$)			Phi-3-mini-4k-instruct ($p = 0.05$) ($p = 0.1$) ($p = 0.2$)		
FullFT	Harmful Rate (R_h)	26.0%	53.7%	93.5%	59.4%	90.7%	98.8%
	Harmful Score (S_h)	1.88	2.97	4.53	3.23	4.45	4.89
SPPFT	Harmful Rate (R_h)	7.9%	25.6%	68.3%	8.1%	23.3%	78.8%
	Harmful Score (S_h)	1.25	1.86	3.61	1.29	1.74	3.88
Lisa	Harmful Rate (R_h)	20.7%	41.4%	80.3%	41.9%	61.3%	87.4%
	Harmful Score (S_h)	1.76	2.38	4.01	2.49	3.10	4.23

Table 3: Under the settings of $p = 0.05$, $p = 0.1$ and $p = 0.2$, the harmful rate and harmful score of the LLMs obtained after SPPFT, FullFT and Lisa (Huang et al., 2024b) were adopted, respectively.

The results demonstrate that the effectiveness of SPPFT in mitigating security degradation from harmful data fine-tuning attacks compared to full-parameter fine-tuning across various malicious data rates (p). Under our experimental setup, Lisa (Huang et al., 2024b) falls short of achieving the same level of security preservation as SPPFT. These findings highlight the strong potential of SPPFT in safeguarding models against such attacks.

The strong performance of SPPFT across these four diverse scenarios highlights its versatility in defending against general fine-tuning attacks.

Ablation Experiment. We also conducted ablation studies to show that, freezing the parameters of contiguous layers other than the safety layers during the fine-tuning process does not preserve the security of the aligned LLM. Detailed experimental data can be found in Appendix A.4.5. After comparing with full fine-tuning, freezing the discrete “security-critical neurons” (Wei et al., 2024) and freezing the continuous non-safety layer parameters in fine-tuning, SPPFT proves to be the most effective parameter-freezing defense against security degradation caused by fine-tuning.

5 CONCLUSION

Our work is the first to reveal the security mechanisms within the layer-wise internal parameters of aligned LLMs, confirming the existence of safety layers and developing generalized methods to accurately identify the range of safety layers across different aligned LLMs. Building on this, we propose a novel fine-tuning method, SPPFT, which preserves the security mechanisms by not updating the gradients of the security layers during the fine-tuning process. As a pioneering paper to expose the security mechanisms of aligned LLMs, we hope our research lays a solid groundwork for advancing the field of harmless AI and future developments in LLM security.

ACKNOWLEDGEMENTS

This research was supported by the National Key R&D Program of China 2021YFB2900103, China National Natural Science Foundation with No. 62441228, “the Fundamental Research Funds for the Central Universities” WK2150110024, Science and Technology Tackling Program of Anhui Province, No. 202423k09020016.

REFERENCES

- Marah Abdin, Sam Ade Jacobs, Ammar Ahmad Awan, Jyoti Aneja, Ahmed Awadallah, Hany Awadalla, Nguyen Bach, Amit Bahree, Arash Bakhtiari, Harkirat Behl, et al. Phi-3 technical report: A highly capable language model locally on your phone. *arXiv preprint arXiv:2404.14219*, 2024.
- Josh Achiam, Steven Adler, Sandhini Agarwal, Lama Ahmad, Ilge Akkaya, Florencia Leoni Aleman, Diogo Almeida, Janko Altschmidt, Sam Altman, Shyamal Anadkat, et al. Gpt-4 technical report. *arXiv preprint arXiv:2303.08774*, 2023.
- alpaca finance. Huggingface. <https://huggingface.co/datasets/gbharti/finance-alpaca>, 2024. Accessed: 2024-05-21.
- Andy Arditi, O Balcells, A Syed, W Gurnee, and N Nanda. Refusal in llms is mediated by a single direction. In *Alignment Forum*, pp. 15, 2024.
- Yuntao Bai, Andy Jones, Kamal Ndousse, Amanda Askell, Anna Chen, Nova DasSarma, Dawn Drain, Stanislav Fort, Deep Ganguli, Tom Henighan, et al. Training a helpful and harmless assistant with reinforcement learning from human feedback. *arXiv preprint arXiv:2204.05862*, 2022.
- Federico Bianchi, Mirac Suzgun, Giuseppe Attanasio, Paul Röttger, Dan Jurafsky, Tatsunori Hashimoto, and James Zou. Safety-tuned llamas: Lessons from improving the safety of large language models that follow instructions. *arXiv preprint arXiv:2309.07875*, 2023.
- Justin Cui, Wei-Lin Chiang, Ion Stoica, and Cho-Jui Hsieh. Or-bench: An over-refusal benchmark for large language models. *arXiv preprint arXiv:2405.20947*, 2024.
- Josef Dai, Xuehai Pan, Ruiyang Sun, Jiaming Ji, Xinbo Xu, Mickel Liu, Yizhou Wang, and Yaodong Yang. Safe rlhf: Safe reinforcement learning from human feedback. *arXiv preprint arXiv:2310.12773*, 2023.
- JUST A FEW TOKENS DEEP. Safety alignment should be made more than just a few tokens deep.
- Dan Hendrycks, Collin Burns, Steven Basart, Andrew Critch, Jerry Li, Dawn Song, and Jacob Steinhardt. Aligning ai with shared human values. *Proceedings of the International Conference on Learning Representations (ICLR)*, 2021a.
- Dan Hendrycks, Collin Burns, Steven Basart, Andy Zou, Mantas Mazeika, Dawn Song, and Jacob Steinhardt. Measuring massive multitask language understanding. *Proceedings of the International Conference on Learning Representations (ICLR)*, 2021b.
- Chia-Yi Hsu, Yu-Lin Tsai, Chih-Hsun Lin, Pin-Yu Chen, Chia-Mu Yu, and Chun-Ying Huang. Safe lora: the silver lining of reducing safety risks when fine-tuning large language models. *arXiv preprint arXiv:2405.16833*, 2024.
- Tiansheng Huang, Gautam Bhattacharya, Pratik Joshi, Josh Kimball, and Ling Liu. Antidote: Post-fine-tuning safety alignment for large language models against harmful fine-tuning. *arXiv preprint arXiv:2408.09600*, 2024a.
- Tiansheng Huang, Sihao Hu, Fatih Ilhan, Selim Furkan Tekin, and Ling Liu. Lazy safety alignment for large language models against harmful fine-tuning. *arXiv preprint arXiv:2405.18641*, 2024b.
- Tiansheng Huang, Sihao Hu, and Ling Liu. Vaccine: Perturbation-aware alignment for large language model. *arXiv preprint arXiv:2402.01109*, 2024c.

- Divyanshu Kumar, Anurakt Kumar, Sahil Agarwal, and Prashanth Harshangi. Increased llm vulnerabilities from fine-tuning and quantization. *arXiv preprint arXiv:2404.04392*, 2024.
- Chak Tou Leong, Yi Cheng, Kaishuai Xu, Jian Wang, Hanlin Wang, and Wenjie Li. No two devils alike: Unveiling distinct mechanisms of fine-tuning attacks. *arXiv preprint arXiv:2405.16229*, 2024.
- Xuechen Li, Tianyi Zhang, Yann Dubois, Rohan Taori, Ishaan Gulrajani, Carlos Guestrin, Percy Liang, and Tatsunori B. Hashimoto. AlpacaEval: An automatic evaluator of instruction-following models. https://github.com/tatsu-lab/alpaca_eval, 5 2023.
- Chin-Yew Lin. Rouge: A package for automatic evaluation of summaries. In *Text summarization branches out*, pp. 74–81, 2004.
- Kaifeng Lyu, Haoyu Zhao, Xinran Gu, Dingli Yu, Anirudh Goyal, and Sanjeev Arora. Keeping llms aligned after fine-tuning: The crucial role of prompt templates. *arXiv preprint arXiv:2402.18540*, 2024.
- AI Meta. Introducing meta llama 3: The most capable openly available llm to date. *Meta AI*, 2024.
- Jishnu Mukhoti, Yarin Gal, Philip HS Torr, and Puneet K Dokania. Fine-tuning can cripple your foundation model; preserving features may be the solution. *arXiv preprint arXiv:2308.13320*, 2023.
- Long Ouyang, Jeff Wu, Xu Jiang, Diogo Almeida, Carroll L. Wainwright, Pamela Mishkin, Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, John Schulman, Jacob Hilton, Fraser Kelton, Luke E. Miller, Maddie Simens, Amanda Askell, Peter Welinder, Paul Francis Christiano, Jan Leike, and Ryan J. Lowe. Training language models to follow instructions with human feedback. *ArXiv*, abs/2203.02155, 2022a.
- Long Ouyang, Jeffrey Wu, Xu Jiang, Diogo Almeida, Carroll Wainwright, Pamela Mishkin, Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, et al. Training language models to follow instructions with human feedback. *Advances in neural information processing systems*, 35: 27730–27744, 2022b.
- ShengYun Peng, Pin-Yu Chen, Matthew Hull, and Duen Horng Chau. Navigating the safety landscape: Measuring risks in finetuning large language models. *arXiv preprint arXiv:2405.17374*, 2024.
- Xiangyu Qi, Yi Zeng, Tinghao Xie, Pin-Yu Chen, Ruoxi Jia, Prateek Mittal, and Peter Henderson. Fine-tuning aligned language models compromises safety, even when users do not intend to! *ArXiv*, abs/2310.03693, 2023.
- Rafael Rafailov, Archit Sharma, Eric Mitchell, Christopher D Manning, Stefano Ermon, and Chelsea Finn. Direct preference optimization: Your language model is secretly a reward model. *Advances in Neural Information Processing Systems*, 36, 2024.
- Domenic Rosati, Jan Wehner, Kai Williams, Łukasz Bartoszcze, David Atanasov, Robie Gonzales, Subhabrata Majumdar, Carsten Maple, Hassan Sajjad, and Frank Rudzicz. Representation noising effectively prevents harmful fine-tuning on llms. *arXiv preprint arXiv:2405.14577*, 2024.
- Paul Röttger, Hannah Rose Kirk, Bertie Vidgen, Giuseppe Attanasio, Federico Bianchi, and Dirk Hovy. Xstest: A test suite for identifying exaggerated safety behaviours in large language models. *arXiv preprint arXiv:2308.01263*, 2023.
- John Schulman, Filip Wolski, Prafulla Dhariwal, Alec Radford, and Oleg Klimov. Proximal policy optimization algorithms. *arXiv preprint arXiv:1707.06347*, 2017.
- TOMIZATION STAGE. Buckle up: Robustifying llms at every customization stage via data curation.
- Rishub Tamirisa, Bhargu Bharathi, Long Phan, Andy Zhou, Alice Gatti, Tarun Suresh, Maxwell Lin, Justin Wang, Rowan Wang, Ron Arel, et al. Tamper-resistant safeguards for open-weight llms. *arXiv preprint arXiv:2408.00761*, 2024.