

- Tiansheng Huang, Sihao Hu, Fatih Ilhan, Selim Furkan Tekin, and Ling Liu. 2024b. [Lisa: Lazy Safety Alignment for Large Language Models against Harmful Fine-tuning Attack](#). *Advances in Neural Information Processing Systems*.
- Tiansheng Huang, Sihao Hu, Fatih Ilhan, Selim Furkan Tekin, and Ling Liu. 2025b. [Booster: Tackling Harmful Fine-tuning for Large Language Models via Attenuating Harmful Perturbation](#). In *The Thirteenth International Conference on Learning Representations*.
- Tiansheng Huang, Sihao Hu, Fatih Ilhan, Selim Furkan Tekin, and Ling Liu. 2025c. [Virus: Harmful Fine-tuning Attack for Large Language Models Bypassing Guardrail Moderation](#). *Preprint*, arXiv:2501.17433.
- Tiansheng Huang, Sihao Hu, Fatih Ilhan, Selim Furkan Tekin, Zachary Yahn, Yichang Xu, and Ling Liu. 2025d. [Safety tax: Safety alignment makes your large reasoning models less reasonable](#). *Preprint*, arXiv:2503.00555.
- Tiansheng Huang, Sihao Hu, and Ling Liu. 2024c. [Vaccine: Perturbation-aware Alignment for Large Language Models against Harmful Fine-tuning Attack](#). *Advances in Neural Information Processing Systems*.
- Yue Huang, Chujie Gao, Siyuan Wu, Haoran Wang, Xiangqi Wang, Yujun Zhou, Yanbo Wang, Jiayi Ye, Jiawen Shi, Qihui Zhang, Yuan Li, Han Bao, Zhaoyi Liu, Tianrui Guan, Dongping Chen, Ruoxi Chen, Kehan Guo, Andy Zou, Bryan Hooi Kuen-Yew, and 47 others. 2025e. [On the Trustworthiness of Generative Foundation Models: Guideline, Assessment, and Perspective](#). *Preprint*, arXiv:2502.14296.
- Yue Huang, Lichao Sun, Haoran Wang, Siyuan Wu, Qihui Zhang, Yuan Li, Chujie Gao, Yixin Huang, Wenhan Lyu, Yixuan Zhang, and 1 others. 2024d. [Position: TrustLLM: Trustworthiness in large language models](#). In *Proceedings of the 41st International Conference on Machine Learning*, volume 235 of *Proceedings of Machine Learning Research*, pages 20166–20270. PMLR.
- Gabriel Ilharco, Marco Tulio Ribeiro, Mitchell Wortsman, Suchin Gururangan, Ludwig Schmidt, Hananeh Hajishirzi, and Ali Farhadi. 2023. [Editing models with task arithmetic](#). In *The Eleventh International Conference on Learning Representations*.
- Samyak Jain, Ekdeep Singh Lubana, Kemal Oksuz, Tom Joy, Philip HS Torr, Amartya Sanyal, and Puneet K Dokania. 2024. [What Makes and Breaks Safety Fine-tuning? Mechanistic Study](#). *Advances in Neural Information Processing Systems*, 37:93406–93478.
- Jiaming Ji, Boyuan Chen, Hantao Lou, Donghai Hong, Borong Zhang, Xuehai Pan, Juntao Dai, and Yaodong Yang. 2024a. [Aligner: Achieving efficient alignment through weak-to-strong correction](#). *Advances in Neural Information Processing Systems*, 37:93406–93478.
- Jiaming Ji, Mickel Liu, Josef Dai, Xuehai Pan, Chi Zhang, Ce Bian, Boyuan Chen, Ruiyang Sun, Yizhou Wang, and Yaodong Yang. 2024b. [BeaverTails: Towards Improved Safety Alignment of LLM via a Human-Preference Dataset](#). *Advances in Neural Information Processing Systems*, 36.
- Jiaming Ji, Kaile Wang, Tianyi Qiu, Boyuan Chen, Jiayi Zhou, Changye Li, Hantao Lou, Josef Dai, Yunhuai Liu, and Yaodong Yang. 2024c. [Language models resist alignment: Evidence from data compression](#). *Preprint*, arXiv:2406.06144.
- Albert Q. Jiang, Alexandre Sablayrolles, Arthur Mensch, Chris Bamford, Devendra Singh Chaplot, Diego de las Casas, Florian Bressand, Gianna Lengyel, Guillaume Lample, Lucile Saulnier, L  lio Renard Lavaud, Marie-Anne Lachaux, Pierre Stock, Teven Le Scao, Thibaut Lavril, Thomas Wang, Timoth  e Lacroix, and William El Sayed. 2023. [Mistral 7b](#). *Preprint*, arXiv:2310.06825.
- Fengqing Jiang, Zhangchen Xu, Yuetai Li, Luyao Niu, Zhen Xiang, Bo Li, Bill Yuchen Lin, and Radha Poovendran. 2025. [SafeChain: Safety of Language Models with Long Chain-of-Thought Reasoning Capabilities](#). *Preprint*, arXiv:2502.12025.
- Mintong Kang and Bo Li. 2025. [R²-Guard: Robust reasoning enabled llm guardrail via knowledge-enhanced logical reasoning](#). In *The Thirteenth International Conference on Learning Representations*.
- Nir Kshetri. 2024. [Navigating EU Regulations: Challenges for U.S. Technology Firms and the Rise of Europe’s Generative AI Ecosystem](#). *Computer*, 57(10):112–117.
- Chak Tou Leong, Yi Cheng, Kaishuai Xu, Jian Wang, Hanlin Wang, and Wenjie Li. 2024. [No two devils alike: Unveiling distinct mechanisms of fine-tuning attacks](#). *Preprint*, arXiv:2405.16229.
- Simon Lermen, Charlie Rogers-Smith, and Jeffrey Ladish. 2024. [LoRA Fine-tuning Efficiently Undoes Safety Training in Llama 2-Chat 70B](#). *Preprint*, arXiv:2310.20624.
- Mingjie Li, Wai Man Si, Michael Backes, Yang Zhang, and Yisen Wang. 2025. [SaLoRA: Safety-Alignment Preserved Low-Rank Adaptation](#). In *The Thirteenth International Conference on Learning Representations*.
- Guozhi Liu, Weiwei Lin, Tiansheng Huang, Ruichao Mo, Qi Mu, and Li Shen. 2025a. [Targeted Vaccine: Safety Alignment for Large Language Models against Harmful Fine-Tuning via Layer-wise Perturbation](#). *Preprint*, arXiv:2410.09760.
- Sijia Liu, Yuanshun Yao, Jinghan Jia, Stephen Casper, Nathalie Baracaldo, Peter Hase, Yuguang Yao, Chris Yuhao Liu, Xiaojun Xu, Hang Li, and 1 others. 2025b. [Rethinking machine unlearning for large language models](#). *Nature Machine Intelligence*, pages 1–14.

- Xiao-Yang Liu, Guoxuan Wang, Hongyang Yang, and Daochen Zha. 2023. [FinGPT: Democratizing Internet-scale Data for Financial Large Language Models](#). *Preprint*, arXiv:2307.10485.
- Xuyuan Liu, Lei Hsiung, Yaoqing Yang, and Yujun Yan. 2025c. [Spectral insights into data-oblivious critical layers in large language models](#). In *Findings of the Association for Computational Linguistics: ACL 2025*, Vienna, Austria. Association for Computational Linguistics.
- Yang Liu, Yuanshun Yao, Jean-Francois Ton, Xiaoying Zhang, Ruocheng Guo, Hao Cheng, Yegor Klochkov, Muhammad Faaiz Taufiq, and Hang Li. 2024a. [Trust-worthy llms: a survey and guideline for evaluating large language models’ alignment](#). *Preprint*, arXiv:2308.05374.
- Yi Liu, Gelei Deng, Zhengzi Xu, Yuekang Li, Yaowen Zheng, Ying Zhang, Lida Zhao, Tianwei Zhang, Kai-long Wang, and Yang Liu. 2024b. [Jailbreaking chatgpt via prompt engineering: An empirical study](#). *Preprint*, arXiv:2305.13860.
- LlamaTeam. 2024. [The llama 3 herd of models](#). *Preprint*, arXiv:2407.21783.
- Ning Lu, Shengcai Liu, Jiahao Wu, Weiyu Chen, Zhirui Zhang, Yew-Soon Ong, Qi Wang, and Ke Tang. 2025. [Safe Delta: Consistently Preserving Safety when Fine-Tuning LLMs on Diverse Datasets](#). In *Forty-second International Conference on Machine Learning*.
- MetaAI. 2023. [Llama 2 - acceptable use policy - meta ai](#). Accessed on May, 2025.
- MetaAI. 2025. [Developer use guide: your resource for building responsibly](#). Accessed on May, 2025.
- Jishnu Mukhoti, Yarin Gal, Philip HS Torr, and Puneet K Dokania. 2024. [Fine-tuning can cripple your foundation model; preserving features may be the solution](#). *Preprint*, arXiv:2308.13320.
- Long Ouyang, Jeffrey Wu, Xu Jiang, Diogo Almeida, Carroll Wainwright, Pamela Mishkin, Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, John Schulman, Jacob Hilton, Fraser Kelton, Luke Miller, Maddie Simens, Amanda Askell, Peter Welinder, Paul F Christiano, Jan Leike, and Ryan Lowe. 2022. [Training language models to follow instructions with human feedback](#). In *Advances in Neural Information Processing Systems*, volume 35, pages 27730–27744. Curran Associates, Inc.
- ShengYun Peng, Pin-Yu Chen, Matthew Hull, and Duen Horng Chau. 2024. [Navigating the safety landscape: Measuring risks in finetuning large language models](#). In *Advances in Neural Information Processing Systems*.
- Xiangyu Qi, Ashwinee Panda, Kaifeng Lyu, Xiao Ma, Subhrajit Roy, Ahmad Beirami, Prateek Mittal, and Peter Henderson. 2025. [Safety alignment should be made more than just a few tokens deep](#). In *The Thirteenth International Conference on Learning Representations*.
- Xiangyu Qi, Yi Zeng, Tinghao Xie, Pin-Yu Chen, Ruoxi Jia, Prateek Mittal, and Peter Henderson. 2024. [Fine-tuning Aligned Language Models Compromises Safety, Even When Users Do Not Intend To!](#) In *The Twelfth International Conference on Learning Representations*.
- QwenTeam. 2023. [Qwen technical report](#). *Preprint*, arXiv:2309.16609.
- Rafael Rafailov, Archit Sharma, Eric Mitchell, Christopher D Manning, Stefano Ermon, and Chelsea Finn. 2024. [Direct Preference Optimization: Your Language Model is Secretly a Reward Model](#). *Advances in Neural Information Processing Systems*, 36.
- Traian Rebedea, Razvan Dinu, Makesh Narsimhan Sreedhar, Christopher Parisien, and Jonathan Cohen. 2023. [NeMo Guardrails: A Toolkit for Controllable and Safe LLM Applications with Programmable Rails](#). In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing: System Demonstrations*, pages 431–445, Singapore. Association for Computational Linguistics.
- Domenic Rosati, Jan Wehner, Kai Williams, Łukasz Bartoszcze, David Atanasov, Robie Gonzales, Subhabrata Majumdar, Carsten Maple, Hassan Sajjad, and Frank Rudzicz. 2024. [Representation Noising: A Defence Mechanism Against Harmful Finetuning](#). *Advances in Neural Information Processing Systems*.
- Han Shen, Pin-Yu Chen, Payel Das, and Tianyi Chen. 2025. [SEAL: Safety-enhanced Aligned LLM Fine-tuning via Bilevel Data Selection](#). In *The Thirteenth International Conference on Learning Representations*.
- Weijia Shi, Anirudh Ajith, Mengzhou Xia, Yangsibo Huang, Daogao Liu, Terra Blevins, Danqi Chen, and Luke Zettlemoyer. 2024. [Detecting pretraining data from large language models](#). In *The Twelfth International Conference on Learning Representations*.
- Karan Singhal, Tao Tu, Juraj Gottweis, Rory Sayres, Ellery Wulczyn, Mohamed Amin, Le Hou, Kevin Clark, Stephen R Pfohl, Heather Cole-Lewis, and 1 others. 2025. [Toward expert-level medical question answering with large language models](#). *Nature Medicine*, pages 1–8.
- Nisan Stiennon, Long Ouyang, Jeffrey Wu, Daniel Ziegler, Ryan Lowe, Chelsea Voss, Alec Radford, Dario Amodei, and Paul F Christiano. 2020. [Learning to summarize with human feedback](#). *Advances in Neural Information Processing Systems*, 33:3008–3021.
- Rohan Taori, Ishaan Gulrajani, Tianyi Zhang, Yann Dubois, Xuechen Li, Carlos Guestrin, Percy Liang, and Tatsunori B. Hashimoto. 2023. [Stanford alpaca: An instruction-following llama model](#). *GitHub Repository*.

- Hugo Touvron, Louis Martin, Kevin Stone, Peter Albert, Amjad Almahairi, Yasmine Babaei, Nikolay Bashlykov, Soumya Batra, Prajjwal Bhargava, Shruti Bhosale, Dan Bikel, Lukas Blecher, Cristian Canton Ferrer, Moya Chen, Guillem Cucurull, David Es-
 iobu, Jude Fernandes, Jeremy Fu, Wenyin Fu, and 49 others. 2023. [Llama 2: Open Foundation and Fine-Tuned Chat Models](#). *Preprint*, arXiv:2307.09288.
- Cheng Wang, Yue Liu, Baolong Bi, Duzhen Zhang, Zhongzhi Li, Junfeng Fang, and Bryan Hooi. 2025. [Safety in Large Reasoning Models: A Survey](#). *Preprint*, arXiv:2504.17704.
- Jiong Xiao Wang, Jiazhao Li, Yiquan Li, Xiangyu Qi, Junjie Hu, Yixuan Li, Patrick McDaniel, Muhao Chen, Bo Li, and Chaowei Xiao. 2024. [BackdoorAlign: Mitigating Fine-tuning based Jailbreak Attack with Backdoor Enhanced Safety Alignment](#). *Advances in Neural Information Processing Systems*.
- Boyi Wei, Kaixuan Huang, Yangsibo Huang, Tinghao Xie, Xiangyu Qi, Mengzhou Xia, Prateek Mittal, Mengdi Wang, and Peter Henderson. 2024. [Assessing the Brittleness of Safety Alignment via Pruning and Low-Rank Modifications](#). In *Proceedings of the 41st International Conference on Machine Learning*, volume 235 of *Proceedings of Machine Learning Research*, pages 52588–52610. PMLR.
- Di Wu, Xin Lu, Yanyan Zhao, and Bing Qin. 2025. [Separate the Wheat from the Chaff: A Post-Hoc Approach to Safety Re-Alignment for Fine-Tuned Language Models](#). *Preprint*, arXiv:2412.11041.
- Xianjun Yang, Xiao Wang, Qi Zhang, Linda Petzold, William Yang Wang, Xun Zhao, and Dahua Lin. 2023. [Shadow alignment: The ease of subverting safely-aligned language models](#). *Preprint*, arXiv:2310.02949.
- Jingwei Yi, Rui Ye, Qisi Chen, Bin Zhu, Siheng Chen, Defu Lian, Guangzhong Sun, Xing Xie, and Fangzhao Wu. 2024. [On the Vulnerability of Safety Alignment in Open-Access LLMs](#). In *Findings of the Association for Computational Linguistics: ACL 2024*, pages 9236–9260, Bangkok, Thailand. Association for Computational Linguistics.
- Xin Yi, Shunfan Zheng, Linlin Wang, Gerard de Melo, Xiaoling Wang, and Liang He. 2025. [NLSR: Neuron-Level Safety Realignment of Large Language Models Against Harmful Fine-Tuning](#). In *Proceedings of the AAAI Conference on Artificial Intelligence*.
- Yi Zeng, Weiyu Sun, Tran Huynh, Dawn Song, Bo Li, and Ruoxi Jia. 2024. [BEEAR: Embedding-based Adversarial Removal of Safety Backdoors in Instruction-tuned Language Models](#). In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, pages 13189–13215, Miami, Florida, USA. Association for Computational Linguistics.
- Weichao Zhang, Ruqing Zhang, Jiafeng Guo, Maarten de Rijke, Yixing Fan, and Xueqi Cheng. 2024. [Pre-training data detection for large language models: A divergence-based calibration method](#). In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, pages 5263–5274, Miami, Florida, USA. Association for Computational Linguistics.
- Yiran Zhao, Wenxuan Zhang, Yuxi Xie, Anirudh Goyal, Kenji Kawaguchi, and Michael Shieh. 2025. [Understanding and enhancing safety mechanisms of LLMs via safety-specific neuron](#). In *The Thirteenth International Conference on Learning Representations*.
- Lianmin Zheng, Wei-Lin Chiang, Ying Sheng, Siyuan Zhuang, Zhanghao Wu, Yonghao Zhuang, Zi Lin, Zhuohan Li, Dacheng Li, Eric P. Xing, Hao Zhang, Joseph E. Gonzalez, and Ion Stoica. 2023. [Judging LLM-as-a-Judge with MT-Bench and Chatbot Arena](#). *Advances in Neural Information Processing Systems*, 36:46595–46623.
- Kaiwen Zhou, Chengzhi Liu, Xuandong Zhao, Shreedhar Jangam, Jayanth Srinivasa, Gaowen Liu, Dawn Song, and Xin Eric Wang. 2025. [The Hidden Risks of Large Reasoning Models: A Safety Assessment of R1](#). *Preprint*, arXiv:2502.12659.
- Minjun Zhu, Linyi Yang, Yifan Wei, Ningyu Zhang, and Yue Zhang. 2024. [Locking Down the Finetuned LLMs Safety](#). *Preprint*, arXiv:2410.10343.